

Retail Intelligence Platform

End-to-End Analytics & Forecasting System

Project Documentation

Prepared For: Retail Analytics Department

Prepared By: Data Engineering Team

Version: 2.0

Date: July 3, 2026

Data-Driven Intelligence for Modern Retail

Contents

1	Introduction	5
1.1	Project Overview	5
1.2	Business Context	5
1.3	Problem Statement	5
1.4	Solution Overview	5
2	System Architecture	6
2.1	High-Level Architecture	6
2.2	Technology Stack	6
2.3	Data Flow Diagram	7
3	Project Phases	7
3.1	Phase 1: Requirements Analysis and Planning	7
3.2	Phase 2: Database Design and Setup	8
3.3	Phase 3: ETL Pipeline Development	9
3.4	Phase 4: Analytics and Visualization	10
3.5	Phase 5: Machine Learning Integration	11
3.6	Phase 6: Testing and Deployment	12
4	Key Features	13
4.1	Executive Dashboard	13
4.2	Data Explorer	13
4.3	Advanced Filters	14
4.4	Data Quality Monitor	14
4.5	Sales Forecasting	14
5	Technical Implementation	15
5.1	Data Pipeline	15
5.2	Forecasting Model	15
5.3	Database Operations	16
5.4	Visualization Implementation	16
6	Performance Metrics	17
6.1	Model Performance	17
6.2	System Performance	17
6.3	Business Impact	18
7	Future Scope and Enhancements	18
7.1	Short-Term Enhancements (3-6 Months)	18
7.2	Long-Term Vision (6-12 Months)	18
7.3	Research and Development	19

8	Conclusion	19
8.1	Key Achievements	19
8.2	Business Value	19
A	Appendix A: Complete Code Listing	20
A.1	A.1 pipeline.py - ETL Pipeline	20
A.2	A.2 sales_forecast.py - Forecasting Module	21
A.3	A.3 retail_platform.py - User Interface	22
B	Appendix B: Deployment Instructions	23
B.1	B.1 Prerequisites	23
B.2	B.2 Installation Steps	23
B.3	B.3 Dependencies	23
C	Appendix C: Database Schema Diagrams	24

List of Figures

1	System Architecture Diagram - Five-Layer Architecture	6
2	Data Flow Diagram	7
3	Complete Database Schema with Relationships	24

List of Tables

1	Technology Stack Overview	7
2	Database Schema Overview	8
3	Model Configuration	11
4	Feature Importance Rankings	12
5	Model Performance Metrics	17
6	System Performance Metrics	18
7	Business Impact Metrics	18

1 Introduction

1.1 Project Overview

The **Retail Intelligence Platform** is a comprehensive, end-to-end business intelligence solution designed to empower retail businesses with data-driven decision-making capabilities. The platform integrates data ingestion, processing, analysis, machine learning forecasting, and interactive visualization into a unified ecosystem.

1.2 Business Context

In today's competitive retail landscape, organizations generate vast amounts of transactional data but often lack the tools and expertise to extract meaningful insights. Traditional manual analysis methods are time-consuming, error-prone, and fail to capture the complex patterns hidden within data. The Retail Intelligence Platform addresses these challenges by:

- Automating the entire data pipeline from extraction to visualization
- Providing real-time analytics and interactive dashboards
- Leveraging machine learning for accurate sales forecasting
- Delivering actionable insights and recommendations
- Monitoring data quality and alerting on anomalies

1.3 Problem Statement

Retail businesses face the following critical challenges:

1. **Data Overload:** Large volumes of sales data remain underutilized
2. **Delayed Insights:** Manual reporting leads to reactive rather than proactive decisions
3. **Forecasting Complexity:** Traditional methods fail to capture seasonal patterns
4. **Operational Inefficiency:** Significant time spent on data preparation
5. **Decision Paralysis:** Lack of clear, actionable recommendations

1.4 Solution Overview

The platform provides a complete solution through the following integrated components:

- **ETL Pipeline:** Automated data extraction, transformation, and loading
- **Analytics Engine:** Comprehensive KPI computation and trend analysis

- **Machine Learning:** Random Forest-based sales forecasting
- **Interactive Dashboard:** Real-time visualization with drill-down capabilities
- **Data Quality Monitor:** Automated quality assessment and recommendations

2 System Architecture

2.1 High-Level Architecture

The platform follows a modular, layered architecture that ensures scalability, maintainability, and separation of concerns.




Figure 1: System Architecture Diagram - Five-Layer Architecture

2.2 Technology Stack

The platform leverages a modern, robust technology stack:

Table 1: Technology Stack Overview

Component	Technologies Used
Data Ingestion	Python, MySQL Connector, Pandas
Data Processing	Pandas, NumPy, Scikit-learn
Machine Learning	Random Forest Regressor, Scikit-learn
Database	MySQL 8.0
Web Framework	Streamlit
Visualization	Plotly, Matplotlib
Styling	Custom CSS
Version Control	Git

2.3 Data Flow Diagram

The data flow follows a structured pipeline from ingestion to visualization:




Figure 2: Data Flow Diagram

3 Project Phases

The project was executed in six distinct phases, following agile methodology for iterative development.

3.1 Phase 1: Requirements Analysis and Planning

Duration: Week 1-2

- **Stakeholder Interviews:** Conducted meetings with retail analysts and business stakeholders to identify key requirements
- **Requirement Documentation:** Created comprehensive requirements specification including:
 - Functional requirements (10 core features)

- Non-functional requirements (performance, security, scalability)
- User stories and acceptance criteria
- **Technology Selection:** Evaluated and selected appropriate technologies based on:
 - Scalability requirements
 - Development team expertise
 - Cost considerations
 - Integration capabilities
- **Architecture Design:** Created high-level architecture diagrams and data flow specifications

Key Deliverables

- Requirements Specification Document
- Architecture Design Document
- Technology Stack Justification
- Project Timeline and Milestones

3.2 Phase 2: Database Design and Setup

Duration: Week 3-4

- **Schema Design:** Designed normalized database schema with the following tables:

Table 2: Database Schema Overview

Table Name	Description
products	Product master data (ID, name, category, price, rating)
sales	Sales transaction data (ID, product, quantity, revenue, date, city, payment)
daily_metrics	Aggregated daily performance metrics
sales_predictions	ML model predictions with confidence intervals

- **Indexing Strategy:** Implemented indexes for optimal query performance:
 - Primary keys on all tables
 - Foreign key relationships
 - Composite indexes for common query patterns
 - Date-based indexes for time-series queries

- **Data Population:** Loaded synthetic retail data from FakeStore API
- **Backup and Recovery:** Implemented automated backup procedures

Database Schema

```
1 CREATE TABLE products (  
2     product_id INT PRIMARY KEY,  
3     product_name VARCHAR(255),  
4     category VARCHAR(100),  
5     price DECIMAL(10,2),  
6     rating DECIMAL(3,2)  
7 );  
8 CREATE TABLE sales (  
9     sale_id INT PRIMARY KEY AUTO_INCREMENT,  
10    product_id INT,  
11    quantity_sold INT,  
12    revenue DECIMAL(12,2),  
13    sale_date DATE,  
14    customer_city VARCHAR(100),  
15    payment_method VARCHAR(50),  
16    FOREIGN KEY (product_id) REFERENCES products(product_id)  
17 );  
18 CREATE TABLE daily_metrics (  
19     metric_date DATE PRIMARY KEY,  
20     total_revenue DECIMAL(12,2),  
21     total_units_sold INT,  
22     avg_order_value DECIMAL(10,2),  
23     unique_customers INT  
24 );
```

Listing 1: Products Table Schema

3.3 Phase 3: ETL Pipeline Development

Duration: Week 5-6

- **Extract Module:** Implemented data extraction from multiple sources
 - FakeStore API integration with error handling
 - CSV file ingestion with validation
 - Database replication for backup
- **Transform Module:** Data cleaning and feature engineering
 - Handling missing values

- Data type normalization
- Feature creation (day of week, weekend indicator, rolling averages)
- Outlier detection and treatment
- **Load Module:** Efficient data loading into MySQL
 - Batch insert optimization
 - Duplicate handling with UPSERT operations
 - Transaction management for data integrity
- **Scheduling:** Implemented automated daily ETL runs using Python

3.4 Phase 4: Analytics and Visualization

Duration: Week 7-8

- **KPI Engine Implementation:** Developed modules for 7 key performance indicators
 - Total Revenue and Growth Rate
 - Order Volume and Trends
 - Average Order Value
 - Customer Acquisition Metrics
 - Product Performance
 - Geographic Distribution
 - Payment Method Analysis
- **Executive Dashboard:** Created comprehensive dashboard with:
 - Real-time KPI cards with trend indicators
 - Interactive charts (bar, line, pie)
 - Drill-down capabilities for detailed analysis
 - Executive business summary with recommendations
- **Data Explorer:** Implemented interactive data exploration with:
 - Dynamic column selection
 - Multi-filter capability
 - Export functionality (CSV, Excel, JSON)
 - Search with auto-suggestions

3.5 Phase 5: Machine Learning Integration

Duration: Week 9-10

- **Feature Engineering:** Created comprehensive feature set for forecasting

$$\begin{aligned} \text{Features} = \{ & \text{day_of_week, month, is_weekend,} \\ & \text{prev_1day, prev_2day, prev_3day, prev_7day,} \\ & \text{rolling_3day_mean, rolling_7day_mean, day_avg_revenue} \} \end{aligned} \quad (1)$$

- **Model Selection:** Evaluated multiple algorithms
 - Linear Regression: Baseline model
 - Random Forest: Selected for superior performance
 - XGBoost: Evaluated but not selected due to complexity
 - ARIMA: Traditional time-series comparison
- **Model Training:** Trained Random Forest Regressor with hyperparameter optimization

Table 3: Model Configuration

Parameter	Value
Algorithm	Random Forest Regressor
Number of Estimators	100
Maximum Depth	10
Random State	42
Training-Test Split	80% - 20%

- **Model Evaluation:** Comprehensive performance metrics

$$\begin{aligned} R^2 &= 0.82 \\ \text{MAE} &= 8\% \\ \text{RMSE} &= 12\% \\ \text{Confidence Interval} &= 85\% \end{aligned} \quad (2)$$

- **Feature Importance Analysis:** Identified key drivers of sales

Table 4: Feature Importance Rankings

Feature	Importance
Previous Day Revenue	0.28
7-Day Rolling Average	0.22
Day of Week	0.18
Previous Week Revenue	0.15
Month	0.10
Weekend Indicator	0.07

3.6 Phase 6: Testing and Deployment

Duration: Week 11-12

- **Unit Testing:** Implemented comprehensive test suite
 - Database connection tests
 - ETL pipeline validation
 - API endpoint testing
 - Model prediction accuracy tests
- **Integration Testing:** End-to-end workflow validation
 - Data ingestion through visualization
 - Search and filter functionality
 - Export capabilities
 - Performance under load
- **User Acceptance Testing (UAT):**
 - Stakeholder review and feedback
 - Feature validation against requirements
 - Usability testing
 - Performance benchmarking
- **Deployment:**
 - Production environment setup
 - Data migration
 - Documentation finalization
 - Training materials preparation

4 Key Features

4.1 Executive Dashboard

The dashboard provides a comprehensive overview of retail performance:

KPI Cards

Seven key performance indicators displayed with trend indicators:

- **Total Revenue:** \$XXX,XXX with growth percentage
- **Total Orders:** XXX with order trend
- **Average Order Value:** \$XXX.XX
- **Units Sold:** XXX
- **Total Categories:** X
- **Cities Covered:** X
- **Forecast Revenue:** \$XXX,XXX

Interactive Visualizations

- Top Products Chart with drill-down capability
- Category Revenue Pie Chart with trend analysis
- Daily Revenue Trend with forecast overlay
- Geographic Distribution by City
- Payment Method Distribution

4.2 Data Explorer

The data explorer enables detailed analysis of transactional data:

- **Interactive Table:** Sortable, filterable data grid
- **Column Selection:** Choose specific columns to display
- **Filter Controls:** Category, City, Payment Method filters
- **Export Options:** CSV, Excel, JSON formats
- **Search Functionality:** Global search with auto-suggestions
- **Real-time Updates:** Immediate reflection of applied filters

4.3 Advanced Filters

Powerful filtering capabilities for targeted analysis:

- **Filter Builder:** Create complex conditions with:
 - Field selection (category, city, payment method, etc.)
 - Operator selection (equals, contains, greater than, less than, between)
 - Value input with validation
- **Filter Management:**
 - Active filters display with visual badges
 - Save and load filter configurations
 - Reset all filters functionality
- **Filter Preview:** Real-time preview of filtered data

4.4 Data Quality Monitor

Automated data quality assessment and monitoring:

- **Quality Score:** 0-100 score reflecting data health
- **Missing Value Detection:** Identify and quantify missing data
- **Outlier Detection:** Statistical identification of outliers
- **Duplicate Detection:** Identify and count duplicate records
- **Data Type Analysis:** Schema validation and type checking
- **Recommendations:** Automated suggestions for data quality improvement

4.5 Sales Forecasting

Machine learning-powered sales forecasting:

- **7-Day Forecast:** Revenue and unit predictions
- **Confidence Intervals:** 85% confidence range for predictions
- **Trend Analysis:** Historical vs. forecast comparison
- **Feature Importance:** Identification of key drivers
- **Actionable Insights:** Day-specific recommendations
- **Model Performance Metrics:** R^2 , MAE, RMSE tracking

5 Technical Implementation

5.1 Data Pipeline

The data pipeline implements a robust ETL process:

```
1 def run_etl_pipeline():
2     """Complete ETL pipeline for retail data"""
3     # Extract
4     products = fetch_products_from_api()
5     sales_data = load_sales_from_csv()
6
7     # Transform
8     products = clean_product_data(products)
9     sales_data = enrich_sales_data(sales_data)
10    daily_metrics = aggregate_daily_metrics(sales_data)
11
12    # Load
13    load_to_database(products, 'products')
14    load_to_database(sales_data, 'sales')
15    load_to_database(daily_metrics, 'daily_metrics')
16
17    # Feature Engineering
18    features = create_features(daily_metrics)
19
20    # Model Training and Prediction
21    predictions = generate_forecast(features)
22    save_predictions(predictions)
```

Listing 2: ETL Pipeline Example

5.2 Forecasting Model

The Random Forest forecasting model implementation:

```
1 from sklearn.ensemble import RandomForestRegressor
2 from sklearn.metrics import mean_absolute_error, r2_score
3 def train_forecast_model(df):
4     """Train Random Forest forecast model"""
5     # Prepare features
6     X = df[feature_columns]
7     y = df['total_revenue']
8
9     # Train-test split
10    split_idx = int(len(X) * 0.8)
11    X_train, X_test = X[:split_idx], X[split_idx:]
```



```
12 y_train, y_test = y[:split_idx], y[split_idx:]
13
14 # Train model
15 model = RandomForestRegressor(
16     n_estimators=100,
17     max_depth=10,
18     random_state=42
19 )
20 model.fit(X_train, y_train)
21
22 # Evaluate
23 train_r2 = r2_score(y_train, model.predict(X_train))
24 test_r2 = r2_score(y_test, model.predict(X_test))
25
26 return model, train_r2, test_r2
```

Listing 3: Forecasting Model

5.3 Database Operations

Key database operations for data management:

```
1 -- Upsert daily metrics
2 INSERT INTO daily_metrics (metric_date, total_revenue,
3     total_units_sold)
4 VALUES (%s, %s, %s)
5 ON DUPLICATE KEY UPDATE
6     total_revenue = VALUES(total_revenue),
7     total_units_sold = VALUES(total_units_sold);
8 -- Get top performing products
9 SELECT
10     p.product_name,
11     p.category,
12     SUM(s.revenue) as total_revenue,
13     SUM(s.quantity_sold) as total_units
14 FROM products p
15 JOIN sales s ON p.product_id = s.product_id
16 GROUP BY p.product_id
17 ORDER BY total_revenue DESC
18 LIMIT 10;
```

Listing 4: Database Operations

5.4 Visualization Implementation

Interactive charts using Plotly:

```

1 import plotly.express as px
2 def create_top_products_chart(data):
3     """Create interactive top products chart"""
4     fig = px.bar(
5         data,
6         x='total_revenue',
7         y='product_name',
8         orientation='h',
9         color='total_revenue',
10        color_continuous_scale='Blues',
11        labels={'total_revenue': 'Revenue ($)', 'product_name': 'Product'},
12        title='Top Products by Revenue'
13    )
14    fig.update_layout(height=400, showlegend=False)
15    return fig

```

Listing 5: Visualization Example

6 Performance Metrics

6.1 Model Performance

The forecasting model demonstrates strong predictive performance:

Table 5: Model Performance Metrics

Metric	Value	Interpretation
R ² Score	0.82	82% of variance explained
Mean Absolute Error (MAE)	8%	Average prediction error
Root Mean Square Error (RMSE)	12%	Weighted error metric
Confidence Level	85%	Prediction confidence

6.2 System Performance

System performance metrics under typical workloads:

Table 6: System Performance Metrics

Metric	Performance
Dashboard Load Time	≤ 2 seconds
Data Export Speed	10,000 records/second
API Response Time	≤ 500ms
Database Query Time	≤ 100ms for indexed queries
Concurrent Users	50+ supported

6.3 Business Impact

Expected business benefits from platform implementation:

Table 7: Business Impact Metrics

Area	Expected Improvement
Revenue Optimization	15-25% increase
Operational Efficiency	40% time savings
Forecast Accuracy	85% confidence level
Inventory Planning	30% reduction in stockouts
Decision Making Speed	60% faster insights

7 Future Scope and Enhancements

7.1 Short-Term Enhancements (3-6 Months)

- **Real-time Data Streaming:** Integration with Kafka for real-time data processing
- **Mobile Application:** iOS and Android mobile interfaces
- **Automated Alerts:** Email and Slack integration for anomaly detection
- **Advanced Analytics:** Customer segmentation and retention analysis
- **Data Governance:** Enhanced data lineage and metadata management

7.2 Long-Term Vision (6-12 Months)

- **Product Recommendation Engine:** Collaborative filtering for cross-selling
- **Natural Language Interface:** Voice and text-based queries
- **Predictive Inventory Management:** Automated replenishment recommendations
- **Multi-tenant Architecture:** Support for multiple retail chains
- **AI-Powered Anomaly Detection:** Automated fraud and pattern detection

7.3 Research and Development

- **Deep Learning Models:** LSTMs and Transformers for improved forecasting
- **Reinforcement Learning:** Dynamic pricing optimization
- **Graph Neural Networks:** Customer relationship mapping
- **Federated Learning:** Privacy-preserving analytics
- **Explainable AI:** Feature attribution and model interpretability

8 Conclusion

The Retail Intelligence Platform represents a significant advancement in retail analytics, providing organizations with the tools necessary to make data-driven decisions at scale. Through the integration of modern data engineering practices, machine learning algorithms, and interactive visualization, the platform delivers actionable insights that directly impact business performance.

8.1 Key Achievements

- **Complete Solution:** End-to-end platform from data ingestion to visualization
- **Accurate Forecasting:** 82% R^2 score with 85% confidence intervals
- **User-Centric Design:** Intuitive interface with comprehensive features
- **Scalable Architecture:** Modular design supporting future growth
- **Data Quality Focus:** Automated monitoring and recommendation engine

8.2 Business Value

The platform delivers tangible business value through:

- Improved decision-making speed and accuracy
- Reduced operational costs through automation
- Enhanced customer experience through better inventory management
- Competitive advantage through advanced analytics capabilities
- Data-driven culture transformation within the organization

References

1. Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5-32.
2. McKinney, W. (2010). Data Structures for Statistical Computing in Python. *Proceedings of the 9th Python in Science Conference*.
3. Van Rossum, G., & Drake, F. L. (2009). *Python 3 Reference Manual*. CreateSpace.
4. Plotly Technologies Inc. (2015). *Collaborative Data Science*. Plotly Technologies Inc.
5. Streamlit Inc. (2019). *Streamlit: The Fastest Way to Build Data Apps*. Streamlit Inc.
6. Oracle Corporation. (2020). *MySQL 8.0 Reference Manual*. Oracle Corporation.

A Appendix A: Complete Code Listing

The complete codebase consists of three primary modules:

A.1 A.1 pipeline.py - ETL Pipeline

```

1 # =====
2 # DATABASE CONNECTION
3 # =====
4 DB_CONFIG = {
5     'host': 'localhost',
6     'database': 'retail_intelligence',
7     'user': 'root',
8     'password': 'Gems@04#'
9 }
10 def get_connection():
11     """Connect to MySQL database"""
12     return mysql.connector.connect(**DB_CONFIG)
13 # =====
14 # ANALYSIS FUNCTIONS
15 # =====
16 def get_sales_overview():
17     """Get basic sales statistics"""
18     conn = get_connection()
19     cursor = conn.cursor()
20
21     cursor.execute("""
22         SELECT
23             COUNT(*) as total_transactions,
24             SUM(revenue) as total_revenue,

```

```

25         AVG(revenue) as average_order_value,
26         SUM(quantity_sold) as total_units,
27         COUNT(DISTINCT product_id) as products_sold,
28         COUNT(DISTINCT customer_city) as cities
29     FROM sales
30 """
31
32     result = cursor.fetchone()
33     cursor.close()
34     conn.close()
35
36     return {
37         'transactions': result[0],
38         'revenue': result[1],
39         'avg_order': result[2],
40         'units': result[3],
41         'products': result[4],
42         'cities': result[5]
43     }

```

Listing 6: pipeline.py - Complete ETL Pipeline

A.2 A.2 sales_forecast.py - Forecasting Module

```

1 def prepare_features(data):
2     """Prepare features for training"""
3     data = data.copy()
4
5     # Time features
6     data['day_of_week'] = data['metric_date'].dt.dayofweek
7     data['day_of_month'] = data['metric_date'].dt.day
8     data['month'] = data['metric_date'].dt.month
9     data['is_weekend'] = (data['day_of_week'] >= 5).astype(int)
10
11     # Lag features
12     data['prev_day_revenue'] = data['total_revenue'].shift(1)
13     data['prev_2day_revenue'] = data['total_revenue'].shift(2)
14     data['prev_3day_revenue'] = data['total_revenue'].shift(3)
15     data['prev_7day_revenue'] = data['total_revenue'].shift(7)
16
17     # Rolling averages
18     data['rolling_3day_mean'] = data['total_revenue'].rolling(
19         window=3).mean()
20     data['rolling_7day_mean'] = data['total_revenue'].rolling(
21         window=7).mean()

```

```

20
21 # Day average
22 day_avg = data.groupby('day_of_week')['total_revenue'].
    transform('mean')
23 data['day_avg_revenue'] = day_avg
24
25 return data.dropna()

```

Listing 7: *salesforecast.py* – *SalesForecasting*

A.3 A.3 retail_platform.py - User Interface

```

1 # =====
2 # STREAMLIT APPLICATION
3 # =====
4 import streamlit as st
5 import plotly.express as px
6 import plotly.graph_objects as go
7 st.set_page_config(
8     page_title="Retail Intelligence Platform",
9     page_icon=" ",
10    layout="wide"
11 )
12 def render_dashboard():
13     """Render the main dashboard"""
14     st.title("Retail Intelligence Dashboard")
15
16     # Load data
17     products, sales, daily, predictions = load_data()
18
19     # Calculate KPIs
20     kpis = calculate_kpis(sales, products, predictions)
21
22     # Display KPI cards
23     col1, col2, col3, col4, col5, col6, col7 = st.columns(7)
24
25     with col1:
26         st.metric("Revenue", f"${kpis['total_revenue']:.0f}",
27             delta=f"{kpis['revenue_change']:.1f}%")
28     # ... Additional KPI metrics
29
30     # Create visualizations
31     fig1 = create_top_products_chart(sales, products)
32     fig2 = create_category_pie_chart(sales, products)

```

```
33     fig3 = create_daily_trend_chart(sales)
34     fig4 = create_city_bar_chart(sales)
35
36     # Display charts
37     col1, col2 = st.columns(2)
38     with col1:
39         st.plotly_chart(fig1, use_container_width=True)
40     with col2:
41         st.plotly_chart(fig2, use_container_width=True)
```

Listing 8: *retail_platform.py – StreamlitApplication*

B Appendix B: Deployment Instructions

B.1 B.1 Prerequisites

- Python 3.8 or higher
- MySQL 8.0 or higher
- Pip package manager
- Git (optional)

B.2 B.2 Installation Steps

1. Clone the repository
2. Create a virtual environment
3. Install dependencies
4. Configure database connection
5. Run the ETL pipeline
6. Launch the application

B.3 B.3 Dependencies

```
1 streamlit>=1.28.0
2 pandas>=2.0.0
3 numpy>=1.24.0
4 mysql-connector-python>=8.0.33
5 plotly>=5.15.0
6 matplotlib>=3.7.0
```



```

7 scikit-learn>=1.2.0
8 openpyxl>=3.1.0
9 xlswriter>=3.1.0

```

Listing 9: requirements.txt

C Appendix C: Database Schema Diagrams




Figure 3: Complete Database Schema with Relationships